

# Introduction to PyTorch

Lecture 11 for 14-763/18-763

Guannan Qu

Oct 7, 2024

# Traditional ML vs Deep Learning


**Traditional ML** suffers from some issues including:

- Not good at handling high dimensional data (e.g. image and texts).
  - For a  $32 \times 32$  image, # of input features is 1024
  - For a paragraph of texts, can be hundreds of words
- Need to do feature extraction (like Fourier Transform) which is difficult


**Deep Learning** is capable of

- Handling high dimensional data (image, texts)
- No need to do feature extraction
  - Feature extraction is done automatically in deep learning.

# What is Deep Learning?





**Recall: linear regression**



**Deep learning replaces the linear model  
with “layers” of linear models with non-linear activation!**

# What is Deep Learning?

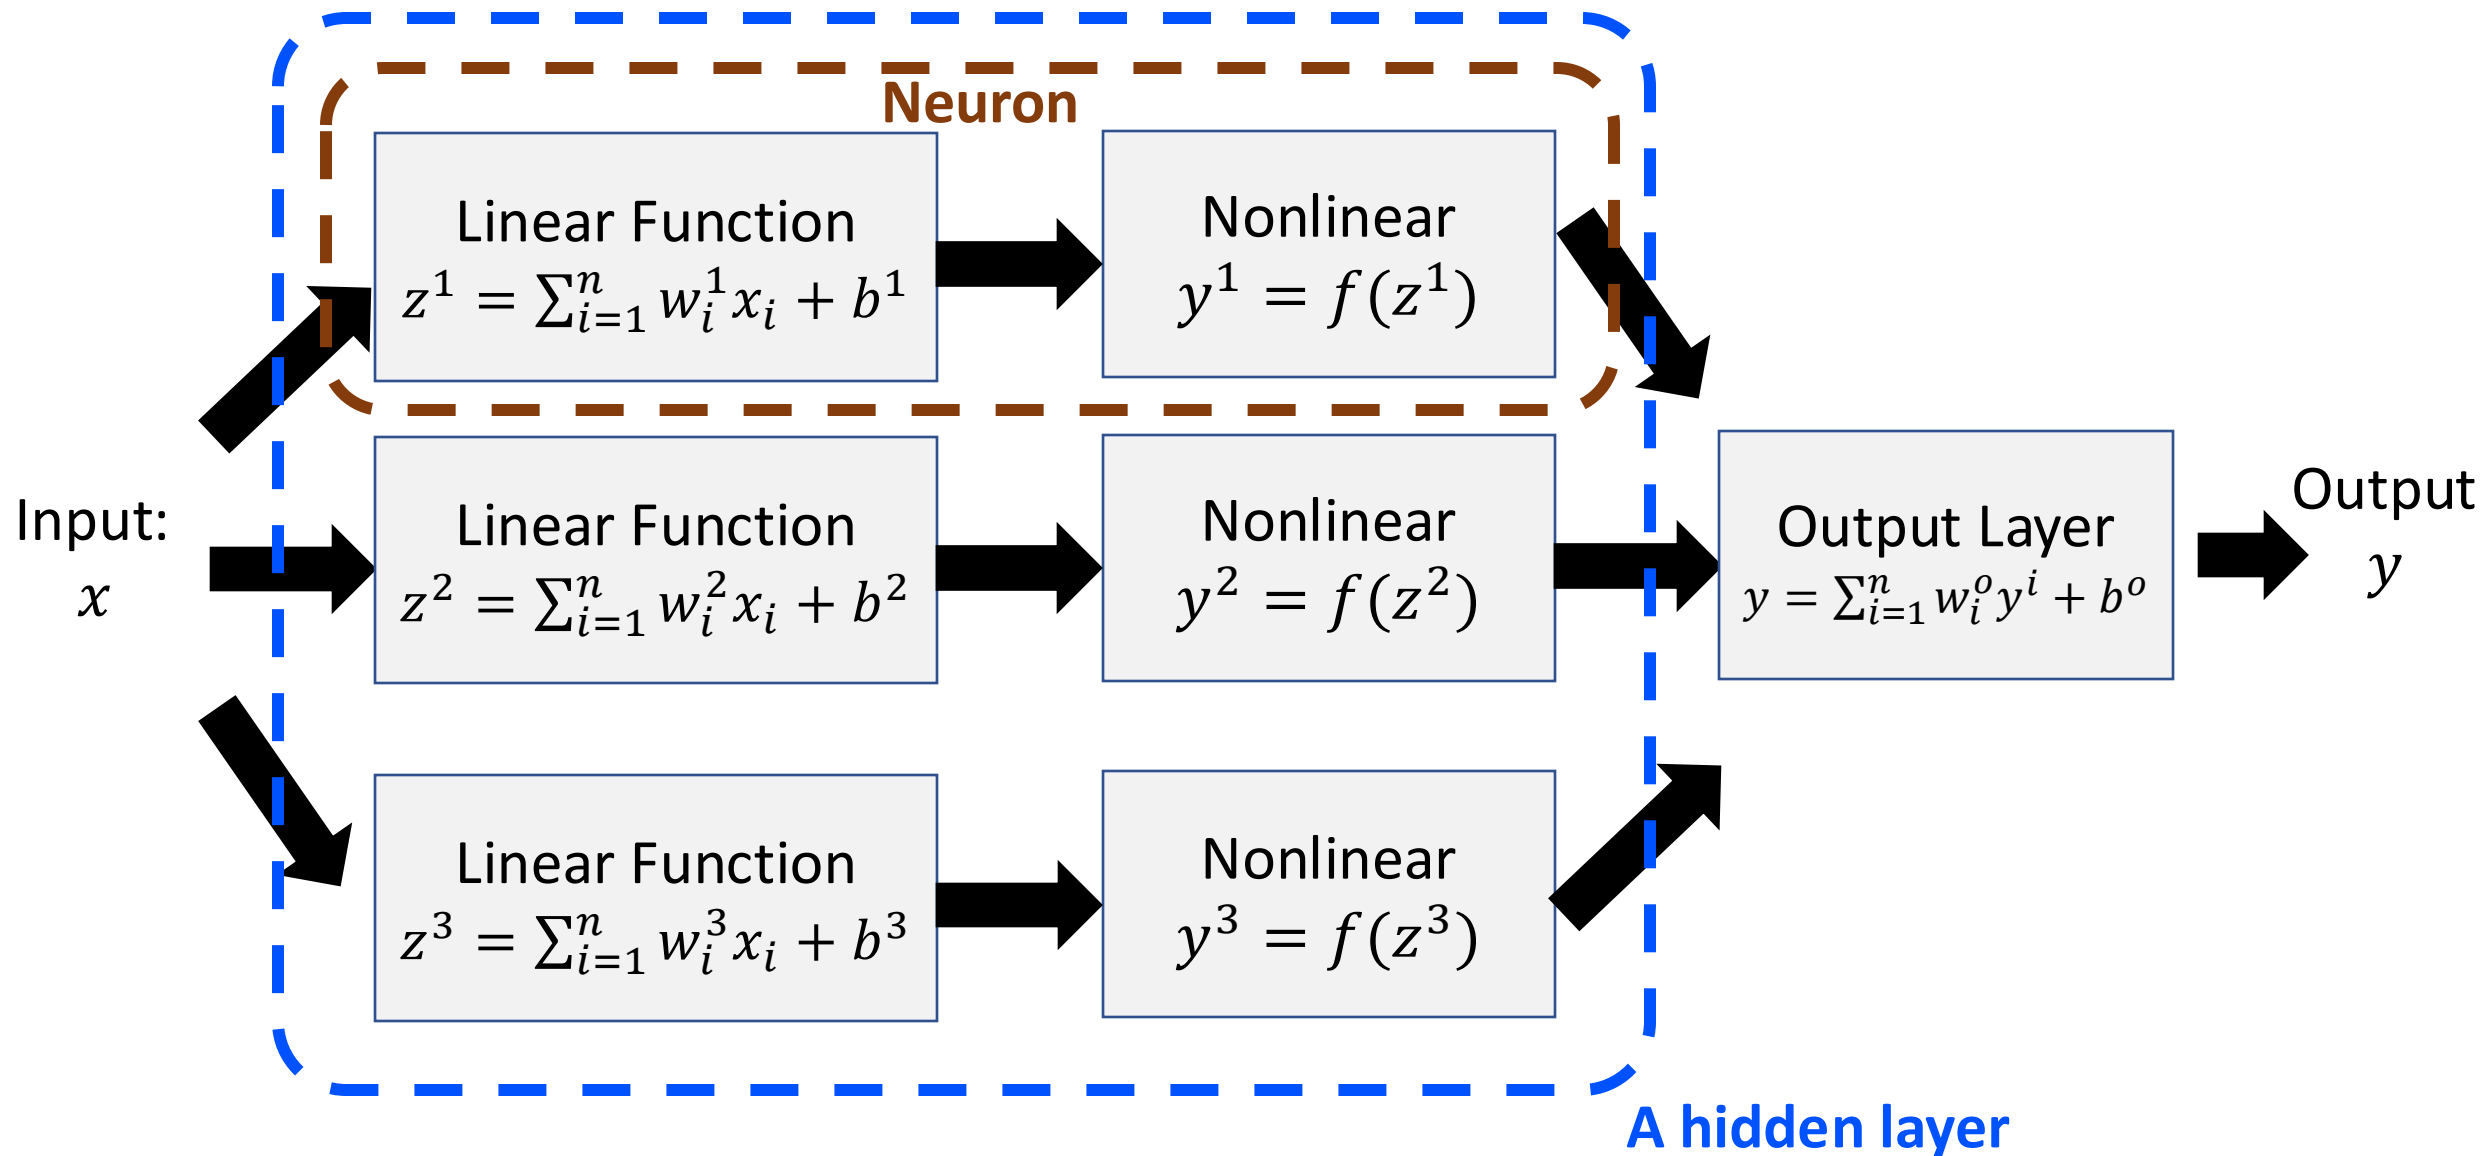
A.k.a. activation function.  
A popular choice is the ReLU  
 $f(y') = \max(y', 0)$



In deep learning, the  
input is typically high-  
dimensional  
 $x = [x_1, \dots, x_n]$


This is a neural network with 1 layer with width 1  
Next: increase the width

# What is Deep Learning?



# What is Deep Learning?

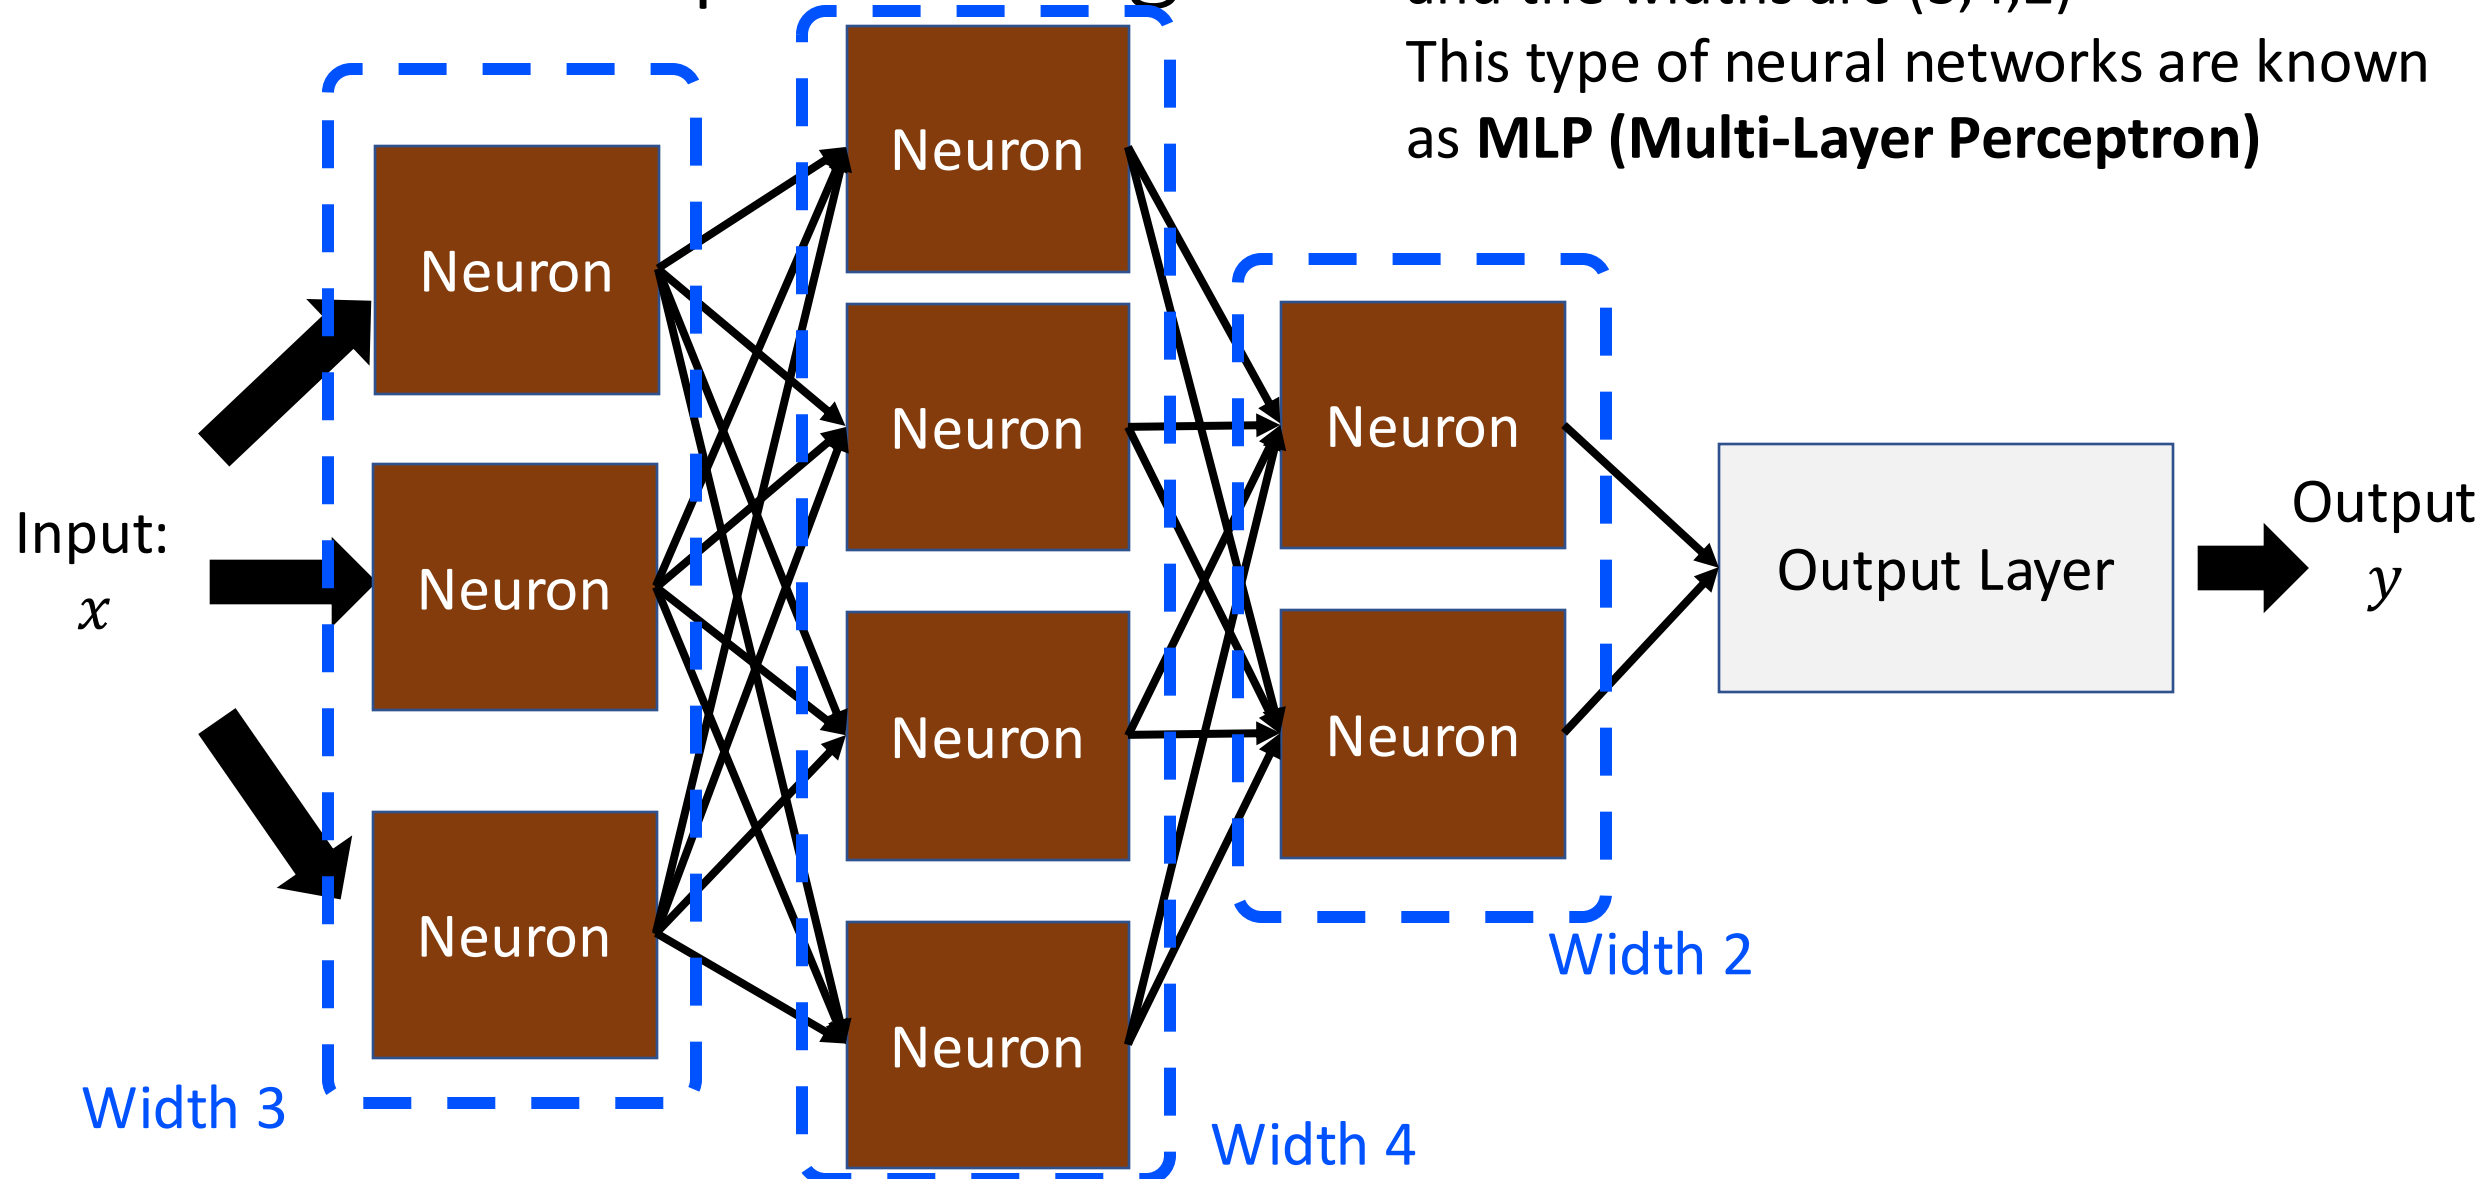
A hidden layer



# What is Deep Learning?

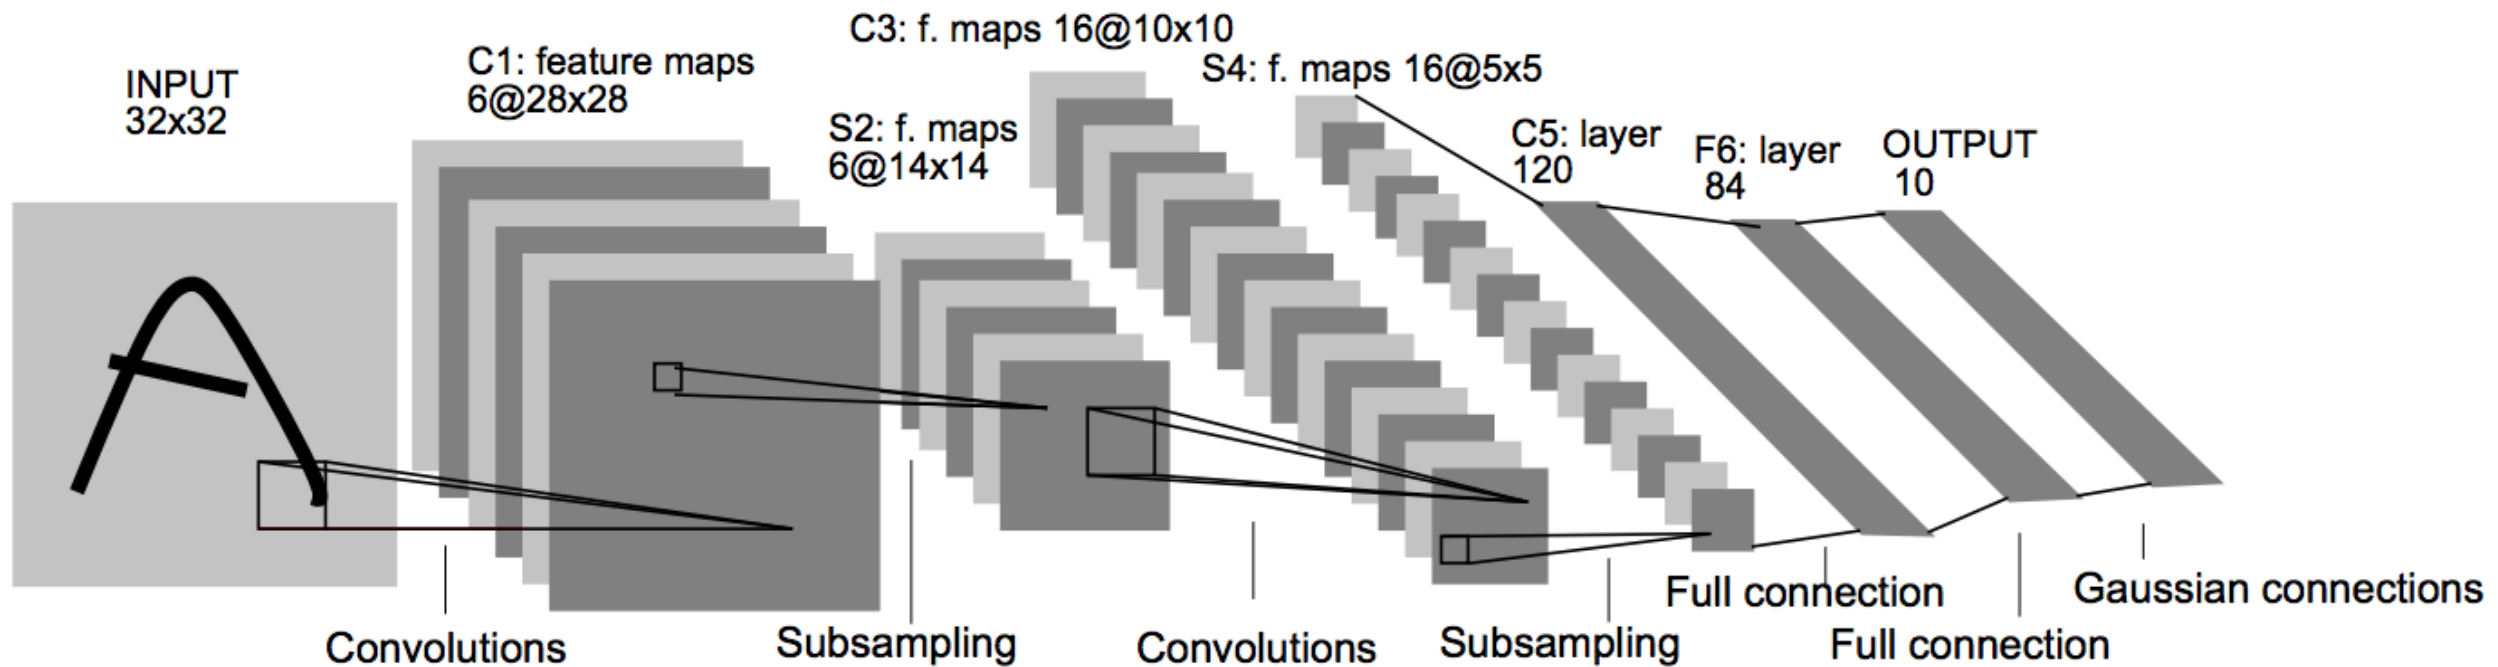
A neural network with 3 hidden layers and the widths are (3,4,2)

This type of neural networks are known as **MLP (Multi-Layer Perceptron)**



# Common NN structures


## LeNet-5 (Convolution NN)





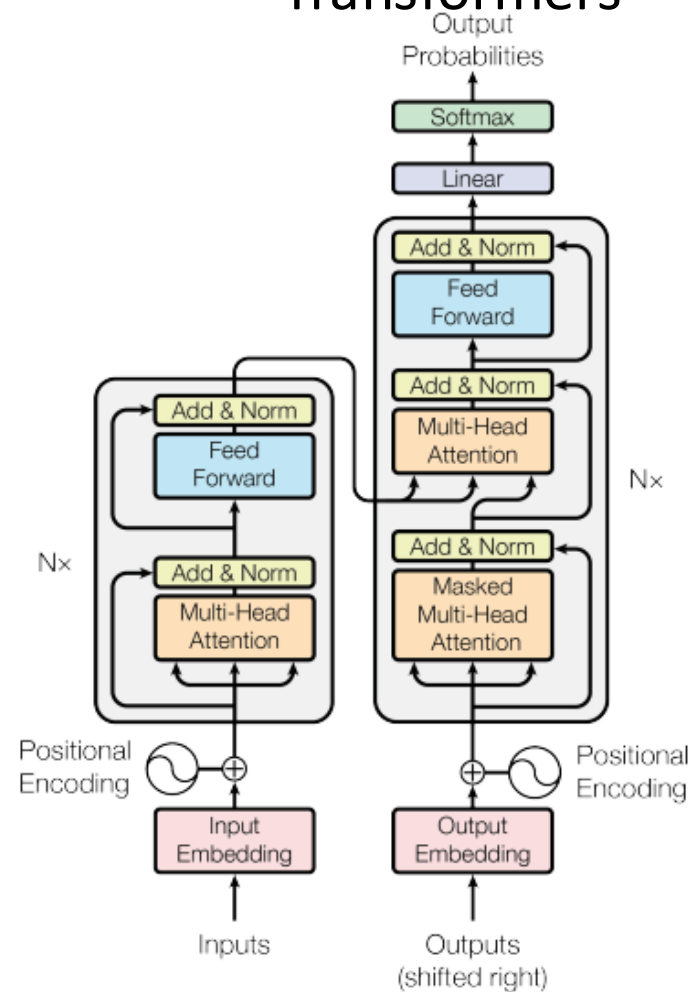
# Common NN structures

LSTM (Recurrent NN)



<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Transformers



<https://arxiv.org/abs/1706.03762>

# What is Deep Learning?

**(Deep) Neural Networks** are a type of ML model:

- Use a cascade of multiple layers of nonlinear processing units (neurons). Each successive layer uses the output from the previous layer as input.
- Has a long history, perhaps first dates back to 1943, but limited success until the 2000s
- Become extremely successful in the 2010s in various domains (image classification, NLP...)
- Various architectures, Multi-Layer Perceptron (MLP), Convolutional NN (CNN), Recurrent NN (RNN), Transformer...

# Deep learning requires new ML platform

SparkML (based on Transformer/Estimator) is not adequate for deep learning

- Deep neural networks has a highly flexible structure
  - # layers, # of neurons for each layer, choice of activation function
  - CNN, RNN, ResNet has more complicated structure
- The training of neural network requires a lot of tuning, and therefore needs to get to the low-level detail
- The training of neural network is data intensive and computationally heavy

**We need specialized ML platform for deep learning!**

# What is TensorFlow (v1)?

- History: Developed by the **Google Brain Team** to accelerate deep neural network research, **TensorFlow v1** was first made public in late 2015
- Built to run on multiple CPUs or GPUs and even mobile operating systems.
- Multiple languages like Python, C/C++ or Java.
- End-to-end, free, and open source. Early dominant player in deep learning.



TensorFlow v1 was the  
early dominant player

2015



# PyTorch

- **PyTorch** was released in 2017 by Facebook AI (now Meta) and soon became popular
  - Known for its simplicity, ease of use, flexibility
  - Uses dynamic computation graph
- In contrast, **TensorFlow v1** at the time
  - Not user friendly, steeper learning curve, not well organized
  - Used static computation graph
  - But TensorFlow still had advantages, e.g. in deployment, visualization



TensorFlow v1 was the  
early dominant player

2015

2017




PyTorch was released  
popular due to ease of use


# TensorFlow v2

The comparison became more complicated when TensorFlow 2 was released in 2019

- TensorFlow 2 became much more user friendly and the APIs were cleaned up
- However, many compatibility issues remained. Code written in TensorFlow v1 cannot easily migrate to v2, frustrating many users
- From 2019 onwards, more and more people (especially in research) switch from tensorflow to PyTorch, but TensorFlow still remain relevant in industry.




# TensorFlow vs PyTorch



<https://www.assemblyai.com/blog/pytorch-vs-tensorflow-in-2023/>

# TensorFlow vs PyTorch




Repository Creation Date

<https://viso.ai/deep-learning/pytorch-vs-tensorflow/>



# TensorFlow vs PyTorch



<https://www.assemblyai.com/blog/pytorch-vs-tensorflow-in-2023/>

# Introduction to PyTorch

Today:

- Use linear regression as warm-up, but go to low level details this time
- Understanding ML: loss function and optimization

This Wed (Oct 9): SGD and Neural Networks

Oct 21: hyper-parameter tuning and best practices

Oct 23: computation graph and GPU

Oct 28: distributed training


Oct 30: ML Ecosystem

# Loss Function

Linear Model:  $y = wx + b$

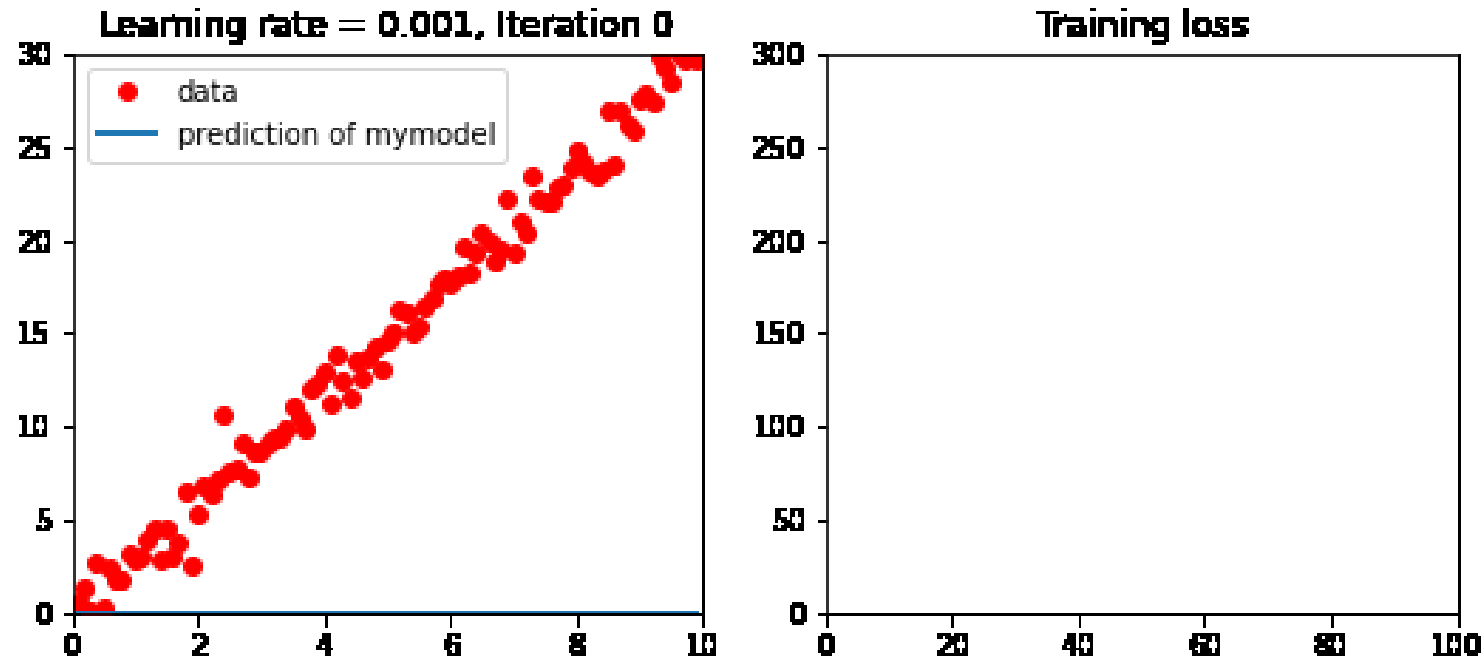
Model Parameters:  $w, b$

$$\text{loss}(w, b) = \frac{1}{N} \sum_i \underbrace{(y_i - (wx_i + b))}_{e_i}^2$$



The training/fitting process finds the  $w, b$  with the smallest loss, but how?

# How does training work?



The training/fitting process finds the  $w, b$  with the smallest loss, but how?

# Optimization: Gradients


Given a function  $loss(\theta)$  that depends on a 2 dimensional  $\theta = [w, b]$

Its gradient  $\nabla loss(\theta)$  is the direction from  $\theta$  that will lead to largest increase in  $loss(\theta)$

# Optimization: Gradients

**Contour Plot of  $loss(\theta)$  (Example)**

Think of it as a terrain altitude map, each circle represents the set of  $\theta$  with the same loss value




Lowest loss achieved at origin

The farther away from origin,  
the larger the loss function

# Optimization: Gradients

To find a parameter with lower loss function, it should follow the **negative gradient direction**

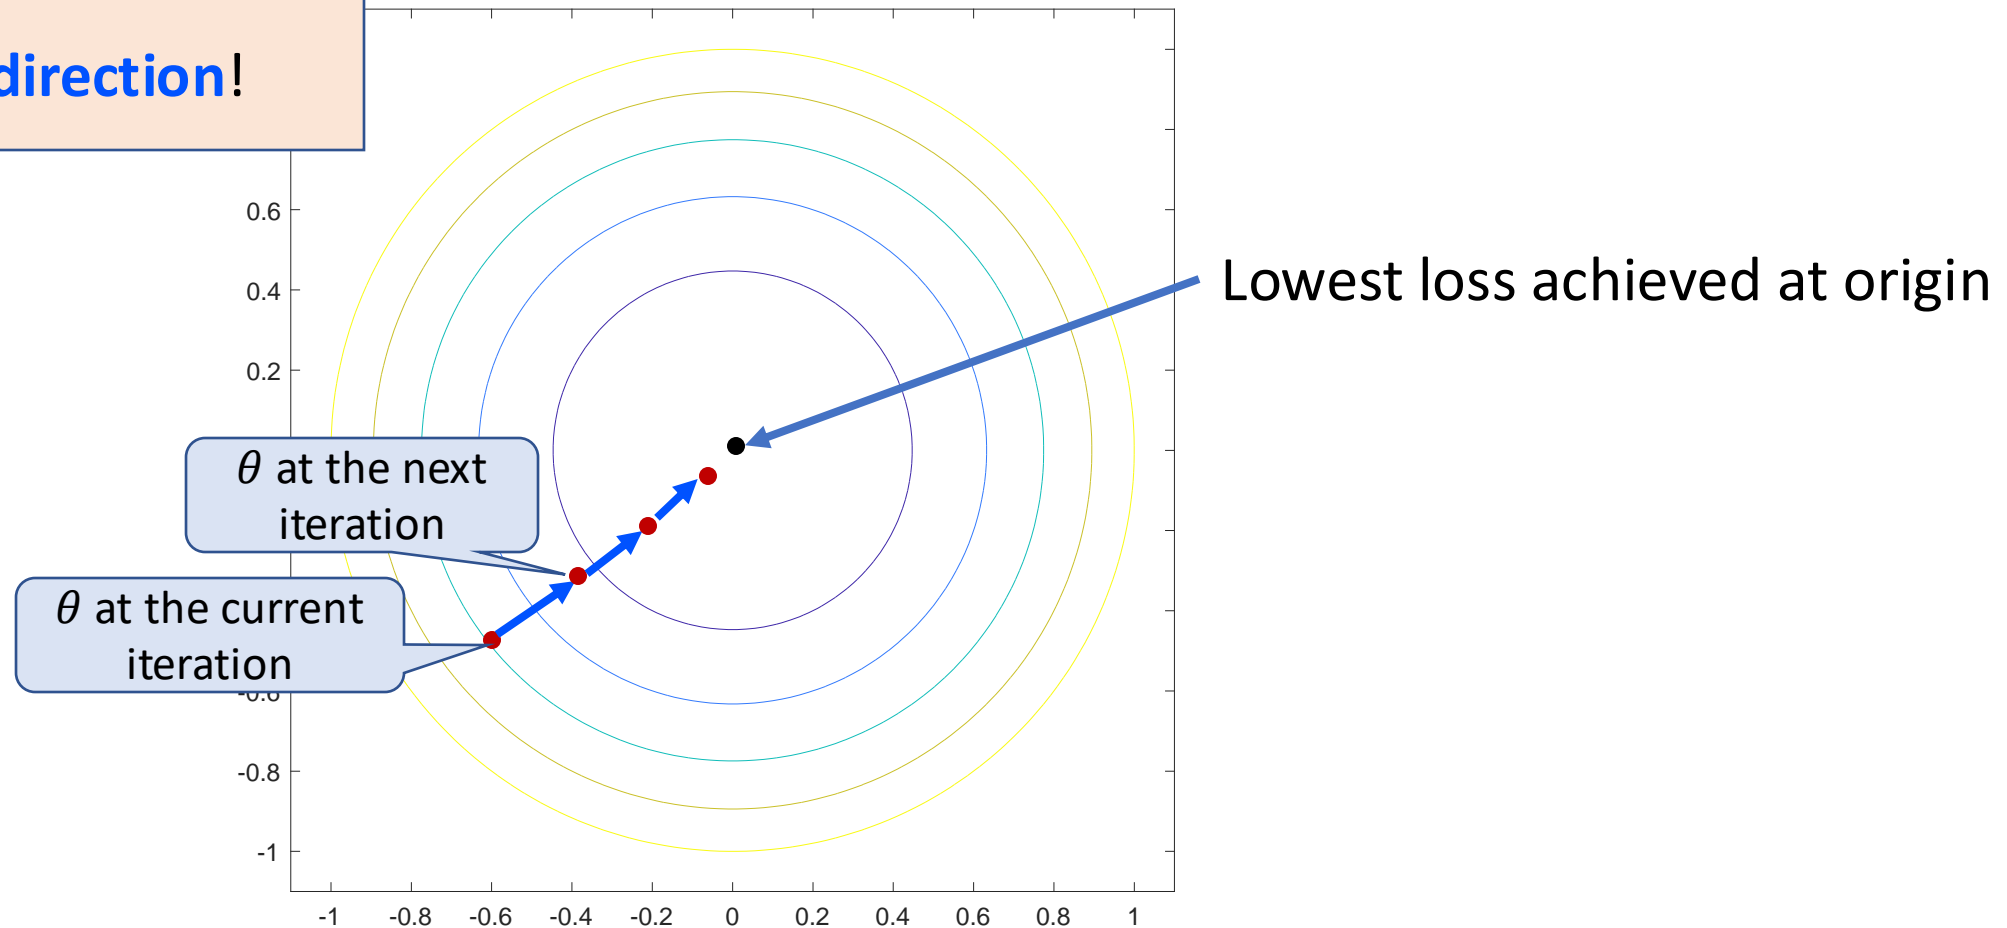


# Optimization: Gradients

## Gradient Descent

Keep following the  
**negative gradient direction!**

Contour Plot of  $loss(\theta)$  (Example)





# Optimization: Gradients

## Gradient Descent


Initialize  $\theta$

Repeat **maxIter** steps:

$$\theta \leftarrow \theta - \eta \nabla \text{loss}(\theta)$$

$\eta$  is learning rate, i.e.  
how large a step one  
makes in each iteration

Contour Plot of  $\text{loss}(\theta)$  (Example)




Lowest loss achieved at origin

# Optimization: Gradients

Linear Model:  $y = wx + b$

Model Parameters:  $w, b$

$$\text{loss}(w, b) = \frac{1}{N} \sum_i (y_i - (wx_i + b))^2$$



**How to calculate the gradient of this loss function?**

# Gradient in PyTorch

The core of PyTorch (and TensorFlow) is their **automatic differentiation (autograd)**

- 1. Define a linear regression model
- 2. Generate some training data
- 3. Calculate gradient and conduct gradient descent

# PyTorch: Linear Regression Model

```
from torch import nn
```

Subclassing nn.Module

```
class MyLinearRegressionModel(nn.Module):
```

```
    def __init__(self, d): # d is the dimension of the input
        super(MyLinearRegressionModel, self).__init__() # call the init function
        # we usually create variables for all our model parameters (w and b in
        # need to create them as nn.Parameter so that the model knows it is an
        self.w = nn.Parameter(torch.zeros(1, d, dtype=torch.float32))
        self.b = nn.Parameter(torch.zeros(1, dtype=torch.float32))
    def forward(self, x):
        # The main purpose of the forward function is to specify given input x,
        return torch.inner(x, self.w) + self.b
```

# PyTorch: Linear Regression Model

In `__init__` function, define all the parameters of the model as `nn.Parameter` and give them initial values

```
class MyLinearRegressionModel(nn.Module):  
    def __init__(self,d): # d is the dimension of the input  
        super(MyLinearRegressionModel,self).__init__() # call the init function  
        # we usually create variables for all our model parameters (w and b in  
        # need to create them as nn.Parameter so that the model knows it is an  
        self.w = nn.Parameter(torch.zeros(1,d, dtype=torch.float32))  
        self.b = nn.Parameter(torch.zeros(1,dtype=torch.float32))
```

Side note:

- In pytorch, tensor is the most basic building block. `nn.Parameter` is a special kind of Tensor used to represent model parameters
- In our code, both parameters are initialized as `torch.zeros`, which are all zero tensors
- Our `__init__` function takes `d` as input, which means the input dimension (we will set `d=1`)

# PyTorch: Linear Regression Model

```
from torch import nn

class MyLinearRegressionModel(nn.Module):
    def __init__(self,d): # d is the dimension of the input
        super(MyLinearRegressionModel,self).__init__() # call the init function
        # we usually create variables for all our model parameters (w and b in
        # need to create them as nn.Parameter so that the model knows it is an
        self.w = nn.Parameter(torch.zeros(1,d, dtype=torch.float32))
        self.b = nn.Parameter(torch.zeros(1,dtype=torch.float32))

    def forward(self,x):
        # The main purpose of the forward function is to specify given input x,
        return torch.inner(x,self.w) + self.b
```

In forward function, define how the output is computed from input. `torch.inner` means inner product, and this line of code simply means  $w_1x_1 + \dots + w_dx_d + b$

# PyTorch: Linear Regression Model

```
mymodel = MyLinearRegressionModel(1) # creating a model instance with input dimension 1
print(mymodel.w)
print(mymodel.b)
```

```
Parameter containing:
tensor([[0.]], requires_grad=True)
Parameter containing:
tensor([0.], requires_grad=True)
```

```
x = torch.tensor(2)
print(mymodel(x)) # we should expect this to be  $w \cdot x + b = 0 \cdot 2 + 0 = 0$  (
```

```
tensor([[0.]], grad_fn=<AddBackward0>)
```



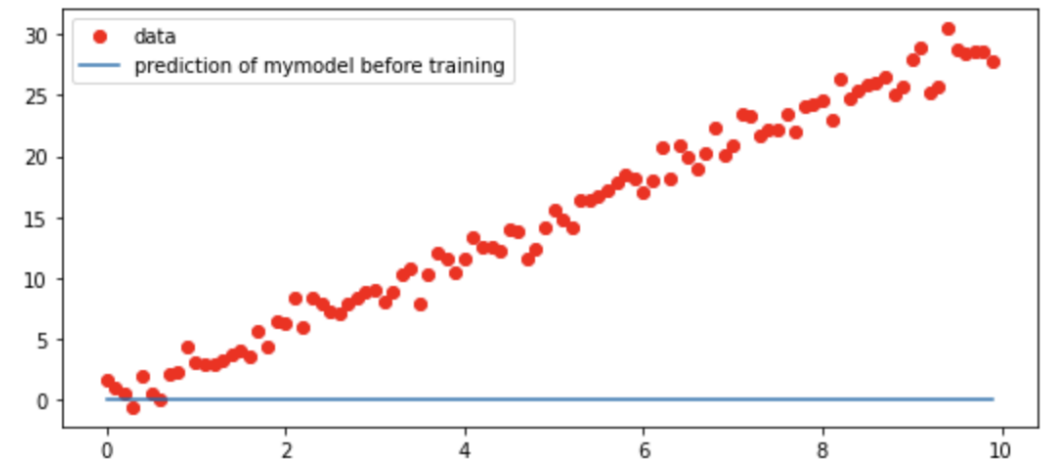
# Gradient in PyTorch

The core of PyTorch (and TensorFlow) is their **automatic differentiation (autograd)**

1. Define a linear regression model
  2. Generate some training data
  3. Calculate gradient and conduct gradient descent
- ← Up Next

```
x = torch.arange(0,10,.1,dtype=torch.float)
x = x[:,None]
y = x*3+torch.randn(x.shape)

prediction = mymodel(x).detach().numpy()
plt.plot(x,y,'ro')
plt.plot(x,prediction)
plt.legend(['data','prediction of mymodel before training'])
```





# Gradient in PyTorch

Recall: we want to calculate the gradient of this loss function

$$loss(w, b) = \frac{1}{N} \sum_i (y_i - (wx_i + b))^2$$

## Steps in PyTorch:

- Step 1: Forward Pass, calculate the loss function value

```
prediction = mymodel(x)

loss = torch.mean((prediction - y)**2)

print(loss)
```

✓ 0.1s

```
tensor(296.1156, grad_fn=<MeanBackward0>)
```

# Gradient in PyTorch

Recall: we want to calculate the gradient of this loss function

$$\text{loss}(w, b) = \frac{1}{N} \sum_i (y_i - (wx_i + b))^2$$

## Steps in PyTorch:

- Step 2: Backward pass.

Before doing backward, let's first check the gradient values now

```
print(my_model.w.grad, my_model.b.grad)
```

✓ 0.6s

None None

# Gradient in PyTorch

Recall: we want to calculate the gradient of this loss function

$$\text{loss}(w, b) = \frac{1}{N} \sum_i (y_i - (wx_i + b))^2$$

## Steps in PyTorch:

- Step 2: Backward pass.

Let's now do backward pass and check gradient again

```
loss.backward()  
print(model.w.grad, model.b.grad)
```

✓ 0.1s

```
tensor([[ -196.8341]]) tensor([ -29.6277])
```

**Up next: gradient descent, i.e. iteratively compute the gradient and conduct gradient descent!**

```
maxIter = 100
```

```
mymodel = MyLinearRegressionModel(1)
```

Tell the optimizer what is the parameters to optimize!

```
# this creates a optimizer, and we tell optimizer we are optimizing the parameters in mymodel
```

```
optimizer = torch.optim.SGD(mymodel.parameters(), lr = 1e-3)
```

Setting learning rate

```
for _ in range(maxIter):
```

Forward Pass

```
# pass input data to get the prediction outputs by the current model
```

```
prediction = mymodel(x)
```

```
# compare prediction and the actual output and compute the loss
```

```
loss = torch.mean((prediction - y)**2)
```

```
# compute the gradient
```

```
optimizer.zero_grad()
```

```
loss.backward()
```

Backward pass and compute gradient.

Note: VERY IMPORTANT to run `optimizer.zero_grad()` to reset grad to zero! Otherwise, the backward will be incorrect.

```
# update parameters
```

```
optimizer.step()
```

Run a gradient descent on the parameters using the computed gradient and the learning rate

```
print(mymodel.w,mymodel.b)

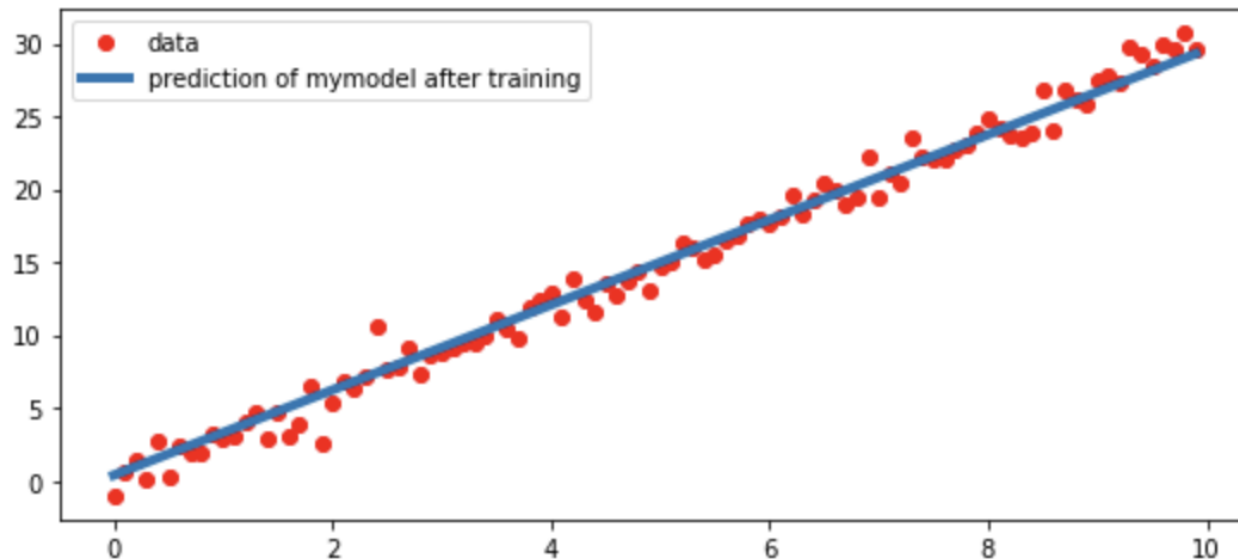
prediction = mymodel(x).detach().numpy()
plt.plot(x,y,'ro')
plt.plot(x,prediction,linewidth = 4)
plt.legend(['data','prediction of mymodel after training'])
```

✓ 0.2s

Parameter containing:

tensor([[2.9180]], requires\_grad=True) Parameter containing:

tensor([0.3993], requires\_grad=True)




# Summary So Far

- Creating a Module: subclassing `nn.Module`
  - Define parameters, define forward function
- Calculating gradient
  - Forward and backward pass
- Perform training (gradient descent)
  - Create optimizer and specify the parameters to optimize, and specify learning rate
  - Write training loop that does gradient descent
    - Forward and compute loss
    - Zero-grad
    - Backward
    - Step

**How to choose learning rate, maxIter? Let's now visualize the gradient descent process!**


# Visualizing Gradient Descent

Learning rate = 0.001



# Visualizing Gradient Descent


Learning rate = 0.0005 (smaller than our first trial)






# Visualizing Gradient Descent

Learning rate = 0.005 (larger than our first trial)



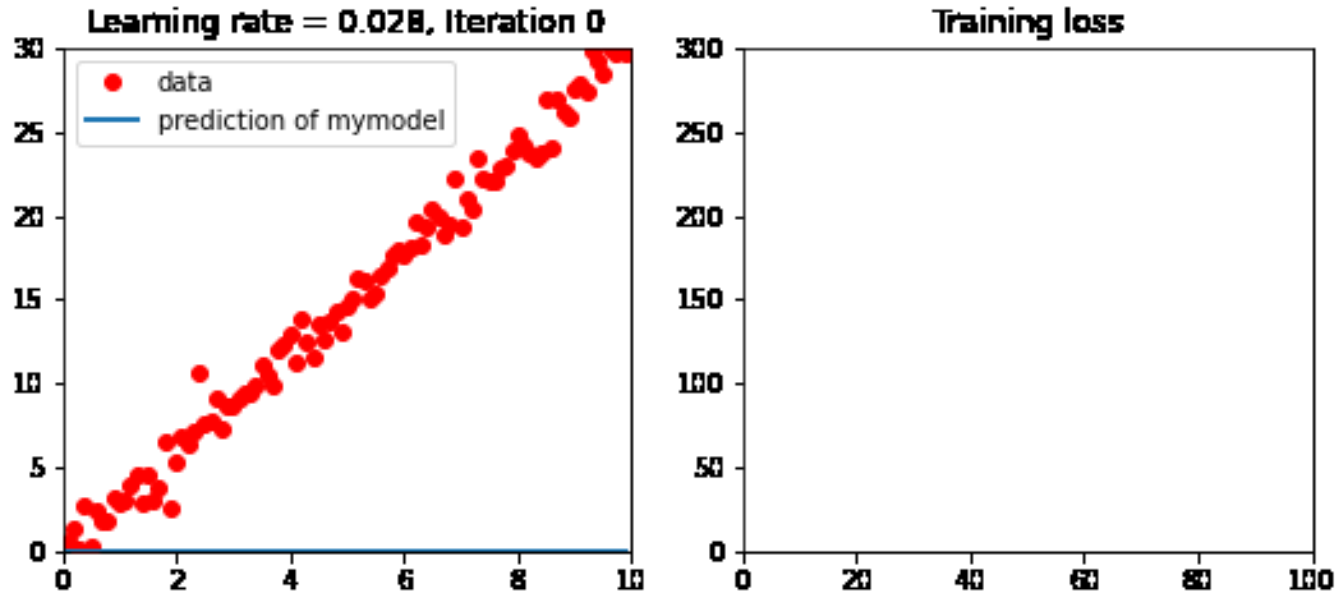
# Visualizing Gradient Descent

Learning rate = 0.025 (much larger than our first trial)




# Visualizing Gradient Descent

Learning rate = 0.028 (much larger than our first trial)



# Visualizing Gradient Descent

Learning rate = 0.03 (much larger than our first trial)



# Lessons Learned on Learning Rate

- Learning rate too small:
  - Converges too slow and takes a lot of iterations
- Learning rate too large:
  - Exhibit unstable (oscillating) behaviors and may diverge
- How to find a good learning rate:
  - Find a small enough learning rate that does not diverge
  - Increase learning rate and plot the training loss curve
    - If the loss curve appears to be converging and “stable”, can further increase
    - If the loss curve appears to be unstable and shows signs of divergence, decrease learning rate