

# PyTorch

## Hyperparameter Tuning and Best Practices

Lecture 13 for 14-763/18-763

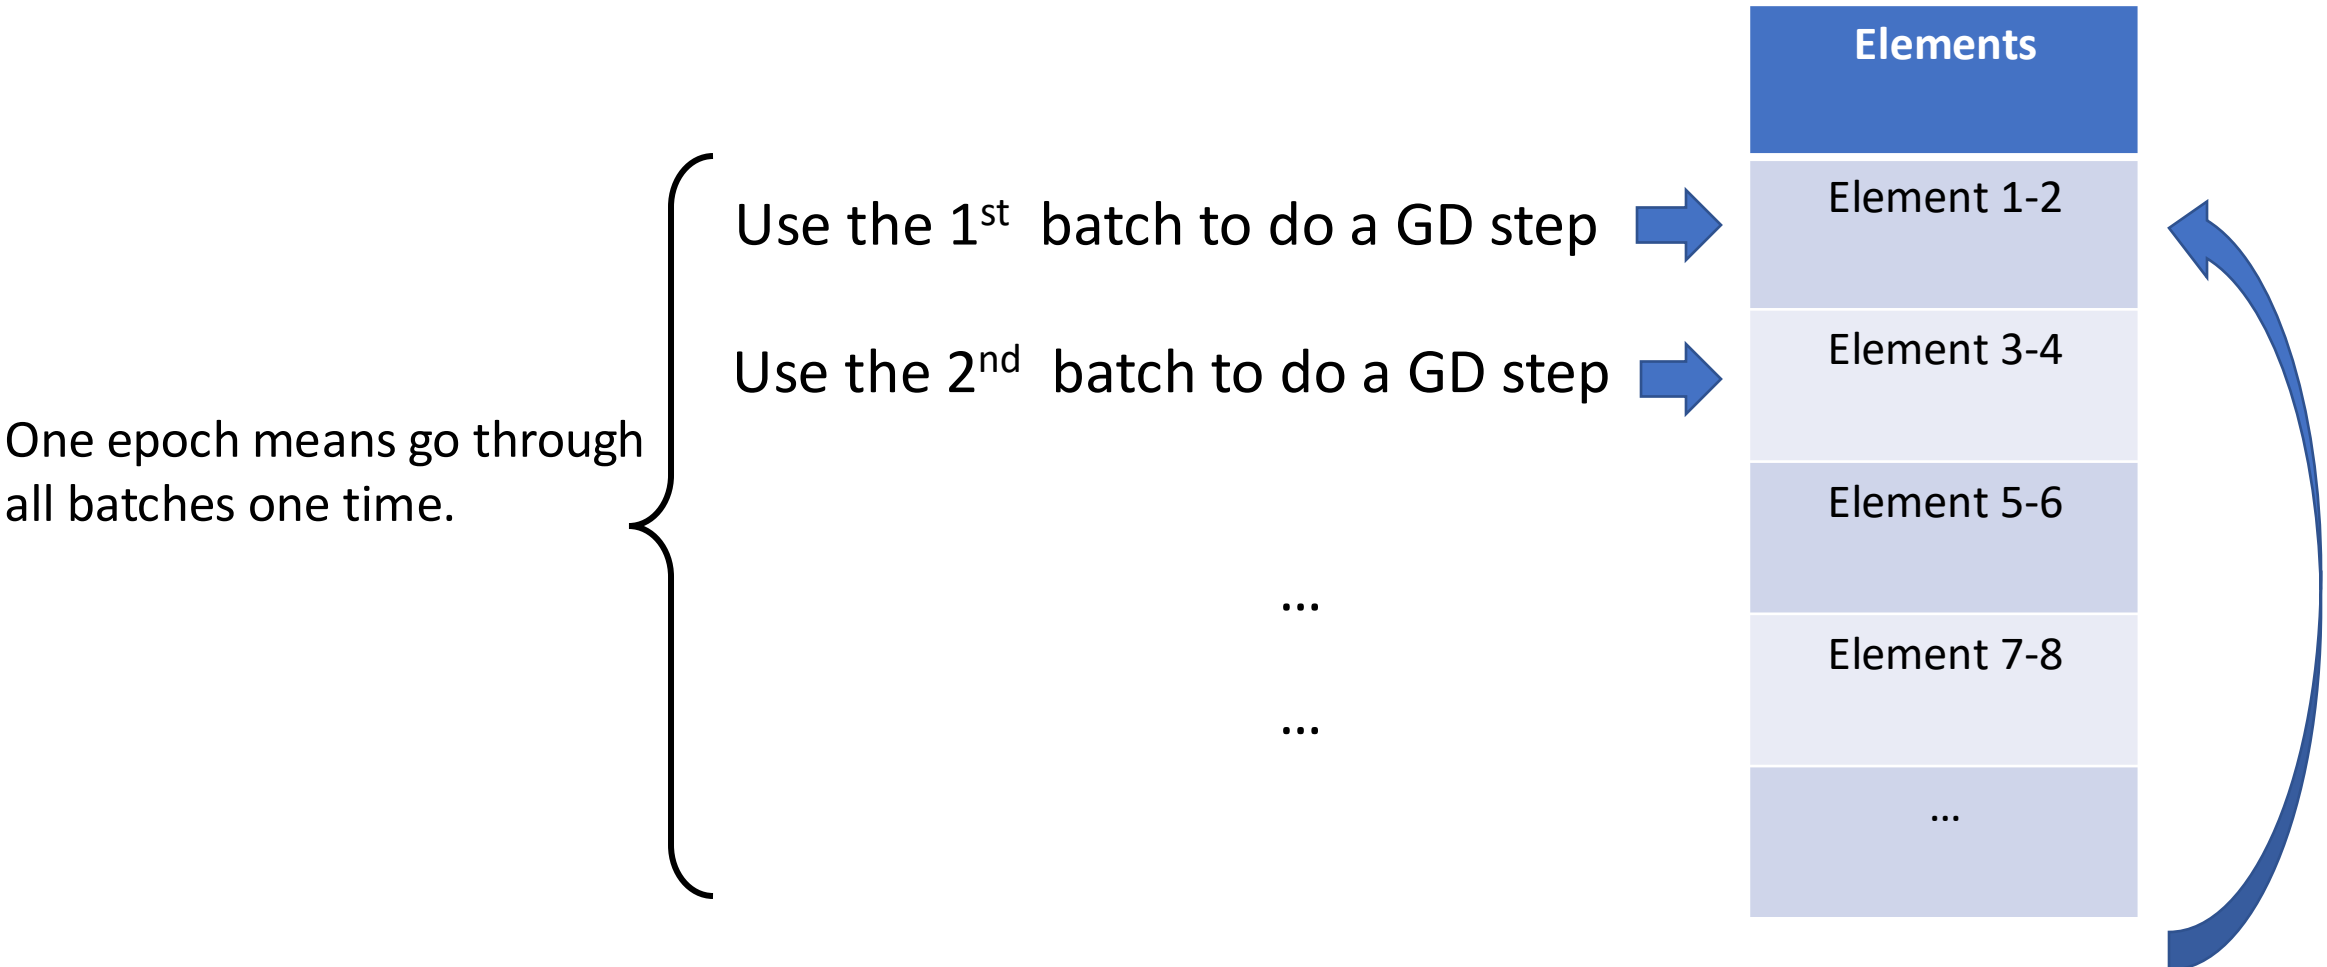
Guannan Qu

Oct 21, 2024

# Recall: Stochastic Gradient Descent


**Batch size = 2**

Batched Dataset




Batch\_size=2  
Learning rate =0.02


Batch size = 2, Learning rate = 0.02, Epoch #0, Batch #0



Training loss




Training loss




Increase batch\_size  
to 8


Batch size = 8, Learning rate = 0.02, Epoch #0, Batch #0



Training loss




Training loss




Batch\_size=2  
Learning rate =0.02


Batch size = 2, Learning rate = 0.02, Epoch #0, Batch #0



Training loss




Training loss




Decrease learning rate  
to 0.001


Batch size = 2, Learning rate = 0.001, Epoch #0, Batch #0



Training loss



Training loss



# Recall: Build Neural Network in PyTorch

```
class myMultiLayerPerceptron(nn.Module):  
    def __init__(self, input_dim, output_dim):  
        super().__init__()
```

Overall, we create a “Sequential” of layers

```
        self.sequential = nn.Sequential( # here we stack multiple layers together
```

```
            nn.Linear(input_dim, 20),  
            nn.ReLU(),
```

The first layer with width 20

```
            nn.Linear(20, 20),  
            nn.ReLU(),
```

The second layer with width 20

```
            nn.Linear(20, 20),  
            nn.ReLU(),
```

The third layer with width 20

```
            nn.Linear(20, 20),  
            nn.ReLU(),
```

The fourth layer with width 20

```
            nn.Linear(20, output_dim)
```

The output layer

```
        )  
    def forward(self, x):  
        y = self.sequential(x)  
        return y
```

# Recall: Training Loops

`mymodel` is now the neural network we just defined

```
mymodel = myMultiLayerPerceptron(1,1) # creating a model instance with input dimension 1
```

```
# Three hyper parameters for training
```

```
lr = .04
```

```
batch_size = 10
```

```
N_epochs = 160
```

Three hyper parameters

```
# Create dataloaders for training and validation
```

```
train_dataloader = DataLoader(train_dataset, batch_size = batch_size, shuffle = True)
```

```
validate_dataloader = DataLoader(validate_dataset, batch_size = batch_size, shuffle = True)
```

```
# Create optimizer
```

```
optimizer = torch.optim.SGD(mymodel.parameters(), lr = lr) # this line creates a optimizer,
```

Dataloader and Optimizer

# Recall: Training Loops

The training loop is identical as before!

```
for epoch in range(N_epochs):  
    batch_loss = []  
    for batch_id, (x_batch, y_batch) in enumerate(train_data_loader):  
        gd_steps+=1  
        # pass input data to get the prediction outputs by the current model  
        prediction = mymodel(x_batch)  
  
        # compare prediction and the actual output and compute the loss  
        loss = torch.mean((prediction - y_batch)**2)  
  
        # compute the gradient  
        optimizer.zero_grad()  
        loss.backward()  
  
        # update parameters  
        optimizer.step()
```

Forward pass

Backward pass and compute gradient.

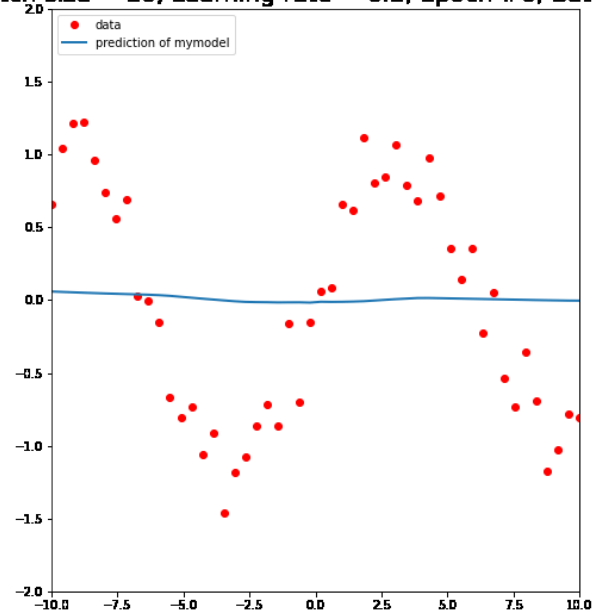
Run a gradient descent step

# Visualizing Neural Network Training


**Let's now visualize the training process and tune hyperparameters!**

Learning rate = 0.2, N\_epochs = 20, batch size = 10


Batch size = 10, Learning rate = 0.2, Epoch #0, Batch #0



Train loss per iteration



Train/validate loss per epoch






# Visualizing Neural Network Training


Let's try a bit more epochs!

Learning rate = 0.2, **N\_epochs = 40**, batch size = 10


Batch size = 10, Learning rate = 0.2, Epoch #0, Batch #0



Train loss per iteration



Train/validate loss per epoch




# Visualizing Neural Network Training


Let's lower the learning rate!

Learning rate = 0.02, N\_epochs = 40, batch size = 10


Batch size = 10, Learning rate = 0.02, Epoch #0, Batch #0



Train loss per iteration



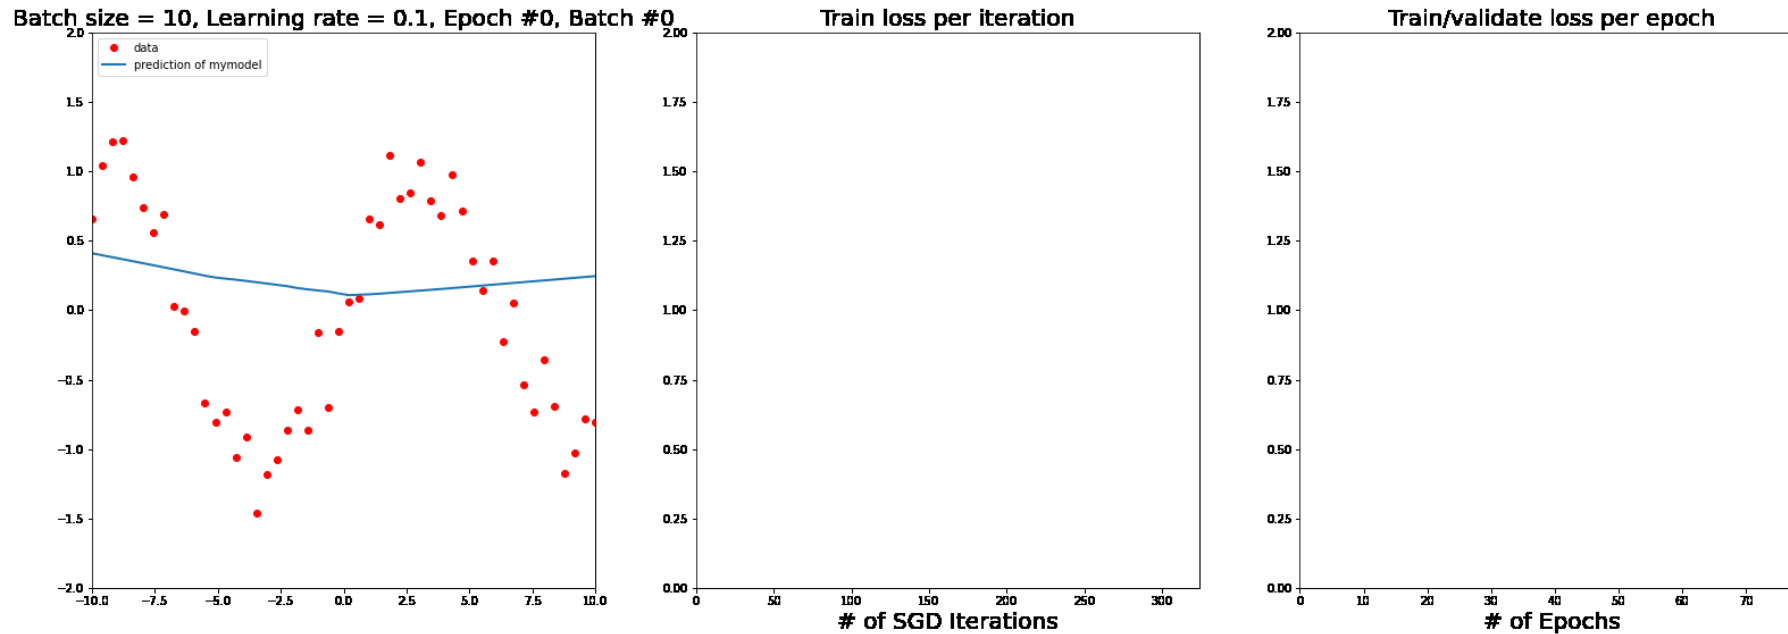
Train/validate loss per epoch



# Visualizing Neural Network Training

Let's increase the learning rate and increase N\_epochs!


Learning rate = 0.1, N\_epochs = 80, batch size = 10



# Visualizing Neural Network Training

Let's lower the learning rate slightly and increase N\_epochs!

Learning rate = 0.04, N\_epochs = 160, batch size = 10



# Lessons Learned on Hyperparameter Tuning

- Too much randomness and oscillations? Loss not improving?
  - Reduce learning rate
  - Increase batch size (this is less common)
- Converging too slow?
  - Increase learning rate
- The loss seems to be decreasing and the model is learning something, but it has not fully converged at the end of the training?
  - Increase the number of epochs

# More on Tuning for Neural Networks

## **Two groups of parameters**

- Training parameters: learning rate, batch\_size, N\_epochs
- Model parameters: depth, width, etc

## **Typical tuning strategy**

- Figure out training parameters that can reliably work well regardless of model parameters
  - This often means being conservative and choose a small learning rate and large number of epochs
  - Fixing the training parameters, try different model parameters

# Summary So Far

Build Models

Subclassing `nn.module`  
Using `nn.Sequential()` to connect `nn.Linear` and `nn.ReLU` together  
Define forward function

Create Dataset

Subclassing `torch.utils.data.Dataset`  
Define `__len__()` and `__getitem__()` function

Training Loops

Create Optimizer, DataLoader  
Nested for-loop, outer-loop for epochs, inner-loop for batches  
Forward, `zero_grad`, backward, step  
Helpful to record train/validate loss (and other metrics) for each epoch

# Up next

- Additional tools
  - Adam optimizer, learning rate scheduler
  - Other built-in activations, loss functions
- NSL-KDD Example: best practices in PyTorch



# Optimizers in PyTorch

Adadelta

Implements Adadelta algorithm.

Adagrad

Implements Adagrad algorithm.

Adam is another very popular optimizer!

Adam

Implements Adam algorithm.

AdamW

Implements AdamW algorithm.

SparseAdam

Implements lazy version of Adam algorithm suitable for sparse tensors.

Adamax

Implements Adamax algorithm (a variation of Adam based on infinity norm).

ASGD

Implements Averaged Stochastic Gradient Descent.

LBFGS

Implements L-BFGS algorithm, heavily inspired by `minFunc`.

NAdam

Implements NAdam algorithm.

RAdam

Implements RAdam algorithm.

RMSprop

Implements RMSprop algorithm.

Rprop

Implements the resilient backpropagation algorithm.

We have been using `torch.optim.SGD`

SGD

Implements stochastic gradient descent (optionally with momentum).

# What is Adam?

Published as a conference paper at ICLR 2015

---

## ADAM: A METHOD FOR STOCHASTIC OPTIMIZATION

**Diederik P. Kingma\***

University of Amsterdam, OpenAI  
dpkingma@openai.com

**Jimmy Lei Ba\***

University of Toronto  
jimmy@psi.utoronto.ca

**Key ingredient** (beyond the scope of this course):

- SGD with momentum
- Rescaled gradients

Very popular, with paper cited more than 190k times

# Implement Adam in PyTorch

```
# Create optimizer – choose between SGD or Adam  
# optimizer = torch.optim.SGD(mymodel.parameters(), lr = lr)  
optimizer = torch.optim.Adam(mymodel.parameters(), lr = lr)
```

Just use `optim.Adam` instead of `optim.SGD`! The rest is identical

# SGD vs Adam

## SGD Optimizer


Batch size = 10

Learning rate = 0.04


N\_epochs = 160

(The parameters we got  
in last lecture)


Batch size = 10, Learning rate = 0.04, Epoch #0, Batch #0



Train loss per iteration



Train/validate loss per epoch



## Adam Optimizer

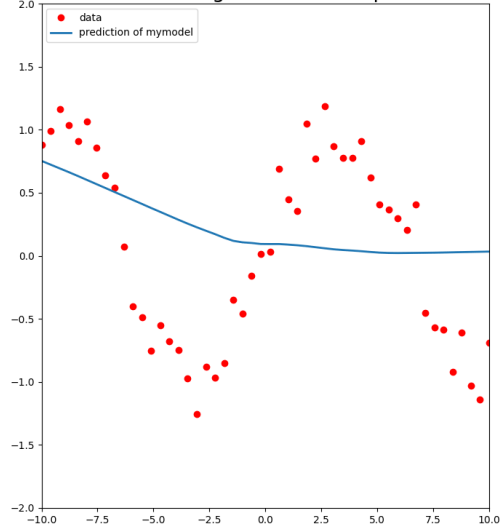
Batch size = 10

Learning rate = 0.02


N\_epochs = 160

(The parameters we got  
in last lecture)


Batch size = 10, Learning rate = 0.02, Epoch #0, Batch #0



Train loss per iteration



Train/validate loss per epoch



# Learning Rate Scheduler

- Recall in the tuning example, we encounter situations where initially the learning is fast, but towards the end the learning is “unstable”
- We had to decrease the learning rate, but a better solution is to use a large learning rate at the beginning, and a smaller learning rate at the end.
- `torch.optim.lr_scheduler` provides methods to adjust learning rate
- For example, `ExponentialLR` shrinks the learning rate by a fixed ratio at each epoch

```
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
scheduler = ExponentialLR(optimizer, gamma=0.9)

for epoch in range(20):
    for input, target in dataset:
        optimizer.zero_grad()
        output = model(input)
        loss = loss_fn(output, target)
        loss.backward()
        optimizer.step()
    scheduler.step()
```

# Built-in Layers in PyTorch

- Containers
- Convolution Layers
- Pooling layers
- Padding Layers
- **Non-linear Activations (weighted sum, nonlinearities)**
- Non-linear Activations (other)
- Normalization Layers
- Recurrent Layers
- Transformer Layers
- Linear Layers
- Dropout Layers
- Sparse Layers
- Distance Functions
- Loss Functions
- Vision Layers
- Shuffle Layers
- DataParallel Layers (multi-GPU, distributed)
- Utilities
- Quantized Functions
- Lazy Modules Initialization

`nn.ReLU`

Applies the rectified linear unit function element-wise:

`nn.ReLU6`

Applies the element-wise function:

`nn.RReLU`

Applies the randomized leaky rectified linear unit function, element-wise, as described in the paper:

`nn.SELU`

Applied element-wise, as:

`nn.CELU`

Applies the element-wise function:

`nn.GELU`

Applies the Gaussian Error Linear Units function:


`nn.Sigmoid`

Applies the element-wise function:

# Comparison of Activation Functions


**Sigmoid**

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




**Leaky ReLU**

$$\max(0.1x, x)$$



**tanh**

$$\tanh(x)$$




**Maxout**

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

**ReLU**

$$\max(0, x)$$



**ELU**

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



# Built-in Layers in PyTorch

- Containers
- Convolution Layers
- Pooling layers
- Padding Layers
- Non-linear Activations (weighted sum, nonlinearity)
- Non-linear Activations (other)
- Normalization Layers
- Recurrent Layers
- Transformer Layers
- Linear Layers
- Dropout Layers
- Sparse Layers
- Distance Functions
- Loss Functions
- Vision Layers
- Shuffle Layers
- DataParallel Layers (multi-GPU, distributed)
- Utilities
- Quantized Functions
- Lazy Modules Initialization

`nn.MSELoss`

Creates a criterion that measures the mean squared error (squared L2 norm) between each element in the input  $x$  and target  $y$ .

`nn.CrossEntropyLoss`

This criterion computes the cross entropy loss between input logits and target.

`nn.KLDivLoss`

The Kullback-Leibler divergence loss.

`nn.BCELoss`

Creates a criterion that measures the Binary Cross Entropy between the target and the input probabilities:



# Example: NSLKDD

- Convert NSL-KDD into `torch.utils.data.Dataset`
- Build NN with **Tanh()** activation
- Training loop with **built-in loss functions**
  - Best practices in writing training loops

# Convert NSL-KDD into torch.utils.data.Dataset

- Convert the spark data frame to pandas, and then numpy
- Use torch.from\_numpy() to convert numpy array to torch tensor

```
x_train = torch.from_numpy(np.array(nslkdd_df_train_pandas['features'].values.tolist(), np.float32))
y_train = torch.from_numpy(np.array(nslkdd_df_train_pandas['outcome'].values.tolist(), np.int64))

x_validate = torch.from_numpy(np.array(nslkdd_df_validate_pandas['features'].values.tolist(), np.float32))
y_validate = torch.from_numpy(np.array(nslkdd_df_validate_pandas['outcome'].values.tolist(), np.int64))

x_test = torch.from_numpy(np.array(nslkdd_df_test_pandas['features'].values.tolist(), np.float32))
y_test = torch.from_numpy(np.array(nslkdd_df_test_pandas['outcome'].values.tolist(), np.int64))
```


# Convert NSL-KDD into torch.utils.data.Dataset

- Convert the tensors to datasets

```
✓ class MyDataset(Dataset):  
✓     def __init__(self,x,y):  
        self.x = x  
        self.y = y  
  
✓     def __len__(self):  
        return self.x.shape[0]  
  
✓     def __getitem__(self, idx):  
        return (self.x[idx],self.y[idx])  
  
train_dataset = MyDataset(x_train,y_train)  
validate_dataset = MyDataset(x_validate,y_validate)  
test_dataset = MyDataset(x_test,y_test)
```

# Neural Network for MultiClass Classification

Neural Network for MultiClass Classification with  $C$  classes



**PyTorch convention:** The neural network's output should be the score (logit) of each class.


**Meaning of score:** the larger the score of a class, the more likely the class is the true output.

The predicted probability of each class

The predicted output is the class that has the **highest probability (the argmax)**

# Neural Network for MultiClass Classification

Neural Network for MultiClass Classification with  $C$  classes

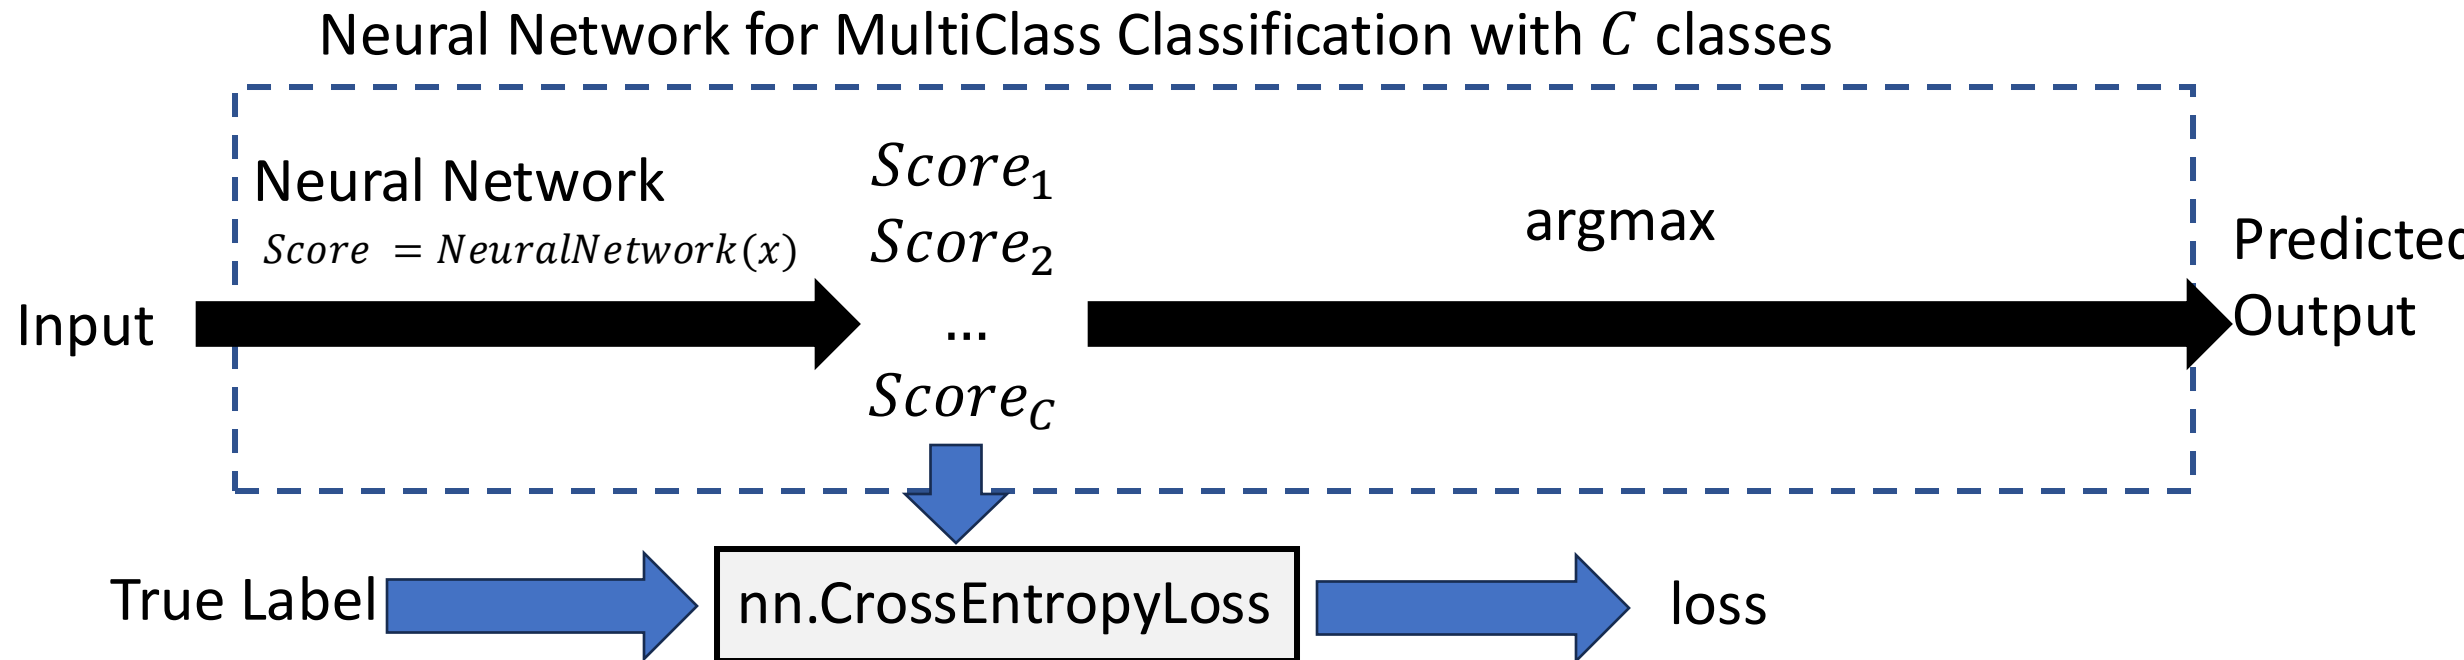


**PyTorch convention:** The neural network's output should be the score (logit) of each class.

**Meaning of score:** the larger the score of a class, the more likely the class is the true output.

Equivalently, the predicted output is the class that has the **highest score** (the **argmax of the scores**)

# Neural Network for MultiClass Classification



Note: we don't use "accuracy" as the loss as accuracy is not differentiable. CrossEntropy is the standard loss for classification problems

```
torch.nn.functional.cross_entropy(input, target, weight=None, size_average=None, ignore_index=-100, reduce=None, reduction='mean', label_smoothing=0.0) [SOURCE]
```

This criterion computes the cross entropy loss between input logits and target.

See [CrossEntropyLoss](#) for details.

## Parameters

- **input** (*Tensor*) – Predicted unnormalized logits; see Shape section below for :

# Build NN with **Tanh()** activation

Suppose input dimension of the NSL-KDD problem is 113, and the output is a label that has two classes (0 or 1). How should we set the output of the neural network?

# Build NN with Tanh() activation

```
class myMultiLayerPerceptron_TahnActivation(nn.Module):
```

```
    def __init__(self, input_dim, output_dim):
```

```
        super().__init__()
```

Overall, we create a “Sequential” of layers

```
        self.sequential = nn.Sequential( # here we stack
```

```
            nn.Linear(input_dim, 20),  
            nn.Tanh(),
```

The first layer with width 20 and Tanh() activation!

```
            nn.Linear(20, 20),  
            nn.Tanh(),
```

The second layer with Tanh() activation!

```
            nn.Linear(20, 20),  
            nn.Tanh(),
```

The third layer with Tanh() activation!

```
            nn.Linear(20, 20),  
            nn.Tanh(),
```

The fourth layer with Tanh() activation!

```
            nn.Linear(20, output_dim)
```

```
        )
```

```
    def forward(self, x):
```

```
        y = self.sequential(x)
```

```
        return y
```

The forward method just uses the sequential we created to compute the output.



# Training Loops: Before the Loop Starts

## Create Model

```
mymodel = myMultiLayerPerceptron_TahnActivation(x_train.shape[1],2) # creating a model instan
```

```
# Three hyper parameters for training
```

```
lr = .005
```

```
batch_size = 64
```

```
N_epochs = 10
```

Define hyper parameters

Good practice: move all hyper parameter definitions to a single place

```
# Create loss function
```

```
loss_fun = nn.CrossEntropyLoss()
```

Create loss function.

```
# Create dataloaders for training and validation
```

```
train_dataloader = DataLoader(train_dataset, batch_size = batch_size, shuffle = True)
```

```
validate_dataloader = DataLoader(validate_dataset, batch_size = batch_size, shuffle = True)
```

```
# Create optimizer
```

```
optimizer = torch.optim.Adam(mymodel.parameters(), lr = lr) # this line creates a optimizer,
```

Create Train/Validate DataLoader and Optimizer

Good practice: set shuffle = True

# Training Loops: Before the Loop Starts

Create lists which we will use to store train/validate loss/accuracy of each epoch

```
losses = [] # training losses of each epoch
accuracies = [] # training accuracies of each epoch

validate_losses = [] # validation losses of each epoch
validate_accuracies = [] # validation accuracies of each epoch

current_best_accuracy = 0.0
```

Keep a variable that will record what is the best validate accuracy seen so far.  
Initialize this to 0

# Training Loops

Create a list that will store the loss/accuracy for each batch

```
for epoch in range(N_epochs):  
    # Train loop  
    batch_loss = [] # keep a list of losses for different batches in this epoch  
    batch_accuracy = [] # keep a list of accuracies for different batches in this epoch  
    for x_batch, y_batch in train_dataloader:
```

.....

(Training loop of this epoch)

← Let's take a look at this

```
    validate_batch_loss = [] # keep a list of losses for different validate batches in this epoch  
    validate_batch_accuracy = []
```

```
    for x_batch, y_batch in validate_dataloader:
```

.....

(Validation loop of this epoch)

Good practice:  
Do validation at each epoch

.....

(End of epoch processing)

# Training Loops: Training

```
batch_loss = [] # keep a list of losses for different batches in this epoch
batch_accuracy = [] # keep a list of accuracies for different batches in this epoch
for x_batch, y_batch in train_dataloader:
```

```
    # pass input data to get the prediction outputs by the current model
```

```
    prediction_score = mymodel(x_batch)
```

```
    # compute the cross entropy loss
```

```
    loss = loss_fun(prediction_score, y_batch)
```

Forward Pass: make predictions and compute loss  
Here we used `loss_fun`, which was created as  
`loss_func = nn.CrossEntropyLoss()`

```
    # compute the gradient
```

```
    optimizer.zero_grad()
```

```
    loss.backward()
```

Zero-grad and backward pass

```
    # update parameters
```

```
    optimizer.step()
```

Optimizer step

```
    # append the loss of this batch to the batch_loss list
```

```
    batch_loss.append(loss.detach().numpy())
```

Record the loss of this batch to the `batch_loss` list

```
    # You can also compute other metrics (accuracy) for this batch here
```

```
    prediction_label = torch.argmax(prediction_score.detach(), dim=1).numpy()
```

```
    batch_accuracy.append(np.sum(prediction_label == y_batch.numpy())/x_batch.shape[0])
```

Calculate accuracy of this batch and record it

# Training Loops: Training

## Details on accuracy calculation:

First need to convert predicted scores to the predicted labels

- The output of NN is `prediction_score`, which is a tensor of shape `(batch_size, 2)`, i.e. 2 entries for each record in the batch
  - The 1<sup>st</sup> entry (index 0) is the score for the class label 0 (normal)
  - The 2<sup>nd</sup> entry (index 1) is the score for the class label 1 (attack)
  - `torch.argmax` will return the index with the highest score, which is the predicted label

After computing the predicted labels, calculate the fraction of them that are correct

```
# You can also compute other metrics (accuracy) for this batch here
prediction_label = torch.argmax(prediction_score.detach(), dim=1).numpy()
batch_accuracy.append( np.sum(prediction_label == y_batch.numpy())/x_batch.shape[0])
```

Calculate accuracy of this batch and record it

# Training Loops: Training

```
for epoch in range(N_epochs):  
    # Train loop  
    batch_loss = [] # keep a list of losses for different batches in this epoch  
    batch_accuracy = [] # keep a list of accuracies for different batches in this epoch  
    for x_batch, y_batch in train_dataloader:
```

.....

(Training loop of this epoch)

```
    validate_batch_loss = [] # keep a list of losses for different validate batches in this epoch  
    validate_batch_accuracy = []
```

```
    for x_batch, y_batch in validate_dataloader:
```

.....

(Validation loop of this epoch)

← Let's take a look at this

.....

(End of epoch processing)

# Training Loops: Validation

```
# Validation loop
validate_batch_loss = [] # keep a list of losses for different validate batches in this epoch
validate_batch_accuracy = []
for x_batch, y_batch in validate_dataloader:
    # pass input data to get the prediction outputs by the current model
    prediction_score = mymodel(x_batch)

    # compare prediction and the actual output and compute the loss
    loss = loss_fun(prediction_score, y_batch)

    # append the lost of this batch to the validate_batch_loss list
    validate_batch_loss.append(loss.detach())

    # You can also compute other metrics (like accuracy) for this batch here
    prediction_label = torch.argmax(prediction_score.detach(), dim=1).numpy()
    validate_batch_accuracy.append(np.sum(prediction_label == y_batch.numpy())/x_batch.shape[0])
```

Forward Pass: make predictions and compute loss

Record the loss of this batch to the batch\_loss list

Calculate accuracy of this batch and record it to list validate\_batch\_loss

# Training Loops

```
for epoch in range(N_epochs):  
    # Train loop  
    batch_loss = [] # keep a list of losses for different batches in this epoch  
    batch_accuracy = [] # keep a list of accuracies for different batches in this epoch  
    for x_batch, y_batch in train_dataloader:
```

.....

(Training loop of this epoch)

```
    validate_batch_loss = [] # keep a list of losses for different validate batches in this epoch  
    validate_batch_accuracy = []
```

```
    for x_batch, y_batch in validate_dataloader:
```

.....

(Validation loop of this epoch)

.....

(End of epoch processing)

← Let's take a look at this



# Training Loops: End of Epoch Processing

```
# calculate the average train loss and validate loss in this epoch and record them
```

```
losses.append(np.mean(np.array(batch_loss)))  
validate_losses.append(np.mean(np.array(validate_batch_loss)))
```

```
# You can also compute other metrics for this epoch here
```

```
accuracies.append(np.mean(np.array(batch_accuracy)))  
validate_accuracies.append(np.mean(np.array(validate_batch_accuracy)))
```

Calculate the average  
train/validate loss across different  
batches in this epoch

Do the same for the accuracy

```
# Printing
```

```
print(f"Epoch = {epoch}, train_loss={losses[-1]}, validate_loss={validate_losses[-1]}")  
print(f"Train accuracy = {np.round(accuracies[-1]*100,2)}%, validate accuracy = {np.round(validate_accuracies[-1]*100,2)}% ")
```


Good practice: Print the progress of this epoch, incl. train/validate loss/accuracy

```
# If the validate metric of this epoch is the best so far, save it  
if validate_accuracies[-1]>current_best_accuracy:  
    print("Current epoch is the best so far. Saving model...")  
    torch.save(mymodel.state_dict(), 'current_best_model')  
    current_best_accuracy = validate_accuracies[-1]
```

Good practice:  
If validate accuracy is the best among all  
previous epochs, save the model

# After Training Loops

Plot the train/validate loss for each epoch



```
fig, axes = plt.subplots(nrows=1, ncols=2)
```

```
axes[0].plot(range(N_epochs), losses, label='Train Loss')
```

```
axes[0].plot(range(N_epochs), validate_losses, label='Validate Loss')
```

```
axes[0].set_xlabel("Epochs")
```

```
axes[0].set_ylabel("Loss")
```

```
axes[1].plot(range(N_epochs), accuracies, label='Train Accuracy')
```

```
axes[1].plot(range(N_epochs), validate_accuracies, label='Validate Accuracy')
```

```
axes[1].set_xlabel("Epochs")
```

```
axes[1].set_ylabel("Accuracy")
```

# After Training Loops

Load the best model encountered in the training loops

```
# create a new model with the same input-output dimension as before
mybestmodel = myMultiLayerPerceptron_TahnActivation(x_train.shape[1],2)

# load the "state_dict" from file into the new model
mybestmodel.load_state_dict(torch.load("current_best_model"))
```

Calculate the test accuracy

```
test_dataloader = DataLoader(test_dataset, batch_size = batch_size, shuffle = True)
test_batch_accuracy = []
for x_batch, y_batch in test_dataloader:
    # pass input data to get the prediction outputs
    prediction_score = mybestmodel(x_batch)

    # Compute metrics (like accuracy) for this batch here
    prediction_label = torch.argmax(prediction_score.detach(), dim=1).numpy()
    test_batch_accuracy.append(np.sum(prediction_label == y_batch.numpy())/x_batch.shape[0])

test_accuracy = np.mean(np.array(test_batch_accuracy))

print(f"Test accuracy = {np.round(test_accuracy*100,2)}%")
```

# Summary of Best Practices

## **Before the training loop begins**

- Define all training hyper-parameters in a single place
- Create (train and validate) DataLoader with shuffle = True

## **Inside the nested training loop**

- Include a validation loop for each epoch
- Record the train/validation loss and metrics
- Print out the progress, incl. epoch, train/validate loss and other metrics
- Save a “best so far” version of the model